

AI Engineering (Week 1)

aie26

Topic 04 : Exploratory Data Analysis

Part 02 : EDA Pass2

Preparation

Dr Bernard Butler

Department of Computing and Mathematics, Waterford Institute of Technology.
(bernard.butler@setu.ie)

Data Handling

Exploring Data 1

Exploring Data 2

Building Models

Autumn Semester, 2026

Outline

- Feature and target distributions
- Plotting
- After EDA Pass 2

Wrap up

AI Engineering (Week 4)

Introduction

Motivating Example

Preparation

Data Handling

Exploring Data 1

Exploring Data 2

Building Models

Prediction

Regression
1

Classification
1

Regression
2

Classification
2

Clustering

Wrap up

Outline

1. Introduction	5
2. A Selection of Statistical Visualisations and Metrics	7
2.1. Categorical Features	8
2.2. Numerical Features	15
3. Summary	21

Acknowledgment

A big thanks to Dr Kieran Murphy, who provided many of the slides for today's lecture.

Introducing EDA Pass2


A Selection of Statistical Visualisations and Metrics

		Relationship to Target	
		Categorical	Numerical
Categorical	Feature	nunique, unique, describe, value_counts, ...	boxplot, ...
		countplot, ...	catplot, boxplot, ...
Numerical	Feature	describe, ...	correlations, ...
		histplot, boxplot, displot, qqplot, ...	lmpplot, ...

Categorical Variables

The Ideal

- Each level equally likely.
- Not too many levels: 2–12(ish).




Tools

- `nunique`, `unique`, `value_counts`.
- `sns.countplot` — shows the counts of observations in each categorical level using bars.

Reality

```
sns.countplot(y="size", data=df);
```

Tips



- If `size` was the target, then most models will train towards the majority class (`size=2`).
- If `size` was a feature, then quality of predictor could vary greatly depending on the feature categorical level.
- Consider merge/drop rare category levels.



Categorical Variables — Relationship with (Categorical) Target

feature

target

Titanic

```
sns.countplot(x="Sex", hue="Survived", data=df);
```



```
pd.crosstab(columns=df.Sex, index=df.Survived, margins=True, normalize='columns') \
    .style.format("{:.2%}").background_gradient(cmap='summer_r')
```

sex	Female	Male	All
smoker			
No	62.07%	61.78%	61.89%
Yes	37.93%	38.22%	38.11%

No relationship between sex and smoker


Sex	female	male	All
Survived			
0	25.80%	81.11%	61.62%
1	74.20%	18.89%	38.38%

Strong relationship between Sex and Survived

Categorical Variables — Relationship with (Numerical) Target

I


```
sns.countplot(x="Size", data=df);
```



- Shows the counts of observations in each categorical level using bar (height/width).

Is it usable?

```
sns.catplot(x="Size", y="a4", data=df, kind='bar');
```




- Shows the average level (mean) and uncertainty (std) of the numerical target (a4) in each categorical level of the categorical variable.
- Vertical bar shows 95% confidence interval.

Is it useful?

Categorical Variables — Relationship with (Numerical) Target




II

The option `kind` in `catplot` can be:



- bar and point show essentially the same information, but point is more compact when comparing multiple categorical features to a continuous target on the same plot.
- strip shows individual observations — useful (as in this case) to show that the larger uncertainty in `Size="small"` observations is mainly due to two outliers.

Example — Dataset: Algae Blooms, Feature: Season, Target: a1






```
df.groupby("Season")["a1"].agg(["mean", "count", "std"])
```

Season	mean count		std
	$\bar{x}$	n	σ
spring	16.649057	53	23.093786
summer	16.038636	44	17.920798
autumn	17.745000	40	21.611203
winter	17.498361	61	22.568256

- Countplot shows no issues with feature **Season** — all levels approximately equally represented.
 - Catplots show slightly less spread in **a1** for **Season="summer"** observations. (**strip** shows smaller range, **point** shows smaller standard deviation).
- ⇒ Mean levels of **a1** for different levels of **Season** are well within the 95% confidence intervals ($\bar{x} \pm \sigma 1.96 / \sqrt{n}$), so no/weak relationship between categorical feature and numerical target.

Example — Dataset: Algae Blooms, Feature: Size, Target: a1



```
df.groupby("Size")["a1"].agg(["min", "max", "mean", "count", "std"])
```

	min	max	mean	count	std
Size			$\bar{x}$	n	σ
small	0.0	89.8	27.255714	70	24.895426
medium	0.0	86.6	11.267857	84	17.163124
large	0.0	56.8	11.611364	44	16.556123

- Countplot shows no issues with feature Size.
 - Catplot (point) shows that levels of a1 are higher for Size="small" observations.
- ⇒ Confidence interval for Size="small" observations do not overlap with CI for other levels, so significant relationship between categorical feature and numerical target.


Numerical Variables

Things here are more complicated as a numerical variable could follow many different distributions. Here we look at data following the standard normal distribution. To start we generate 10,000 values and put in to new DataFrame, df2.

```
rv = stats.norm()
data = rv.rvs(size=10_000)
df2 = pd.DataFrame(data, columns=["x"])
df2.head(5)
```


	x
0	1.011514
1	0.932074
2	0.570263
3	0.001033
4	-0.513810

```
sns.histplot(x="x", data=df2, kde=True);
```



```
sns.boxplot(x="x", data=df2);
```


Histplot (Histogram) and Boxplot



- Histogram is useful in depicting location, spread and shape.
- Curve, is estimate of shape given infinite data and infinite number of bins.
- Boxplots also depicts location, spread and shape, but uses median for estimate of centre, and quartiles for spread.
- Half the data is within the box, data points outside the whiskers (lines) are possible outliers, denoted by circles.


Cumulative Plot and QQ-Plot

```
sns.ecdfplot(data=df2, x="x");
```





- Represents the proportion of observations less than or equal to given value.

```
import pingouin as pg  
pg.qqplot(df2.x);
```




- Plot of observed quantiles against theoretical (assumed normal) quantiles. If both sets of quantiles came from the same distribution, the points lie on a line (approx.).

Example — Dataset: Tips, Feature: total_bill





```
df.total_bill.describe()
```

```
count    244.000000
mean      19.785943
std       8.902412
min       3.070000
25%      13.347500
50%      17.795000
75%      24.127500
max      50.810000
Name: total_bill, dtype: float64
```




- Data is bell curve shaped, but right skewed (data is more spread out to the right).
- Outliers to the right.
- Q-Q-Plot indicate that data is not normal, but we could transform it to be more closer to normal.

Example — Dataset: Titanic, Feature: Fare



```
df.Fare.describe()
```

```
count    891.000000  
mean      32.204208  
std       49.693429  
min        0.000000  
25%       7.910400  
50%      14.454200  
75%      31.000000  
max      512.329200  
Name: Fare, dtype: float64
```




- This variable is more skewed and dominated by its outliers which need to be resolved.


Warning — Plot Output Depends on Data Assumptions

```
df = pd.read_csv("data/train.csv")  
sns.catplot(data=df, x="Fare", y="Survived");
```

```
df = pd.read_csv("data/train.csv")  
df.Survived = df.Survived.astype(str)  
sns.catplot(data=df, x="Fare", y="Survived");
```



- **seaborn** tries to infer the correct graph based on the data values/type, but it does not always get it correct.
- **Survived** stores 0 and 1 and has dtype **int**.
- Converting to a **Categorical** with numeric levels is not enough.
- **astype(str)** converts **0** and **1** to **"0"** and **"1"**.



Summary

- After reading in the data, exploratory data analysis begins
- Pass 1 is all about assessing the structure and cleanliness of the data
 - ① Are the column names descriptive and short, or do we need to rename them?
 - ② What datatype is each column - are there any surprises there?
 - ③ How are missing values handled, and can we standardise this?
- Pass 2 examines the data more closely, in a repeatable fashion
 - ① How are (subsets of) the features and target(s) distributed?
- Pandas and seaborn offer easy-to-use ways of visualising columns, noting
 - ① their datatype
 - ② their cardinality
 - ③ the visualisation objective: observe distributions or relationships within a column